

SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics

**Enhancing Uniping for Distributed
Architecture: Designing a Scalable,
Available, and Fault-Tolerant Black Box
Monitoring System**

Jih-Wei Liang

SCHOOL OF COMPUTATION,
INFORMATION AND TECHNOLOGY —
INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Master's Thesis in Informatics

**Enhancing Uniping for Distributed
Architecture: Designing a Scalable,
Available, and Fault-Tolerant Black Box
Monitoring System**

**Verbesserung von Uniping für verteilte
Architektur: Entwurf eines skalierbaren,
verfügbaren und fehlertoleranten
Black-Box-Überwachungssystems**

Author: Jih-Wei Liang
Supervisor: Bhatotia Pramod; Prof. Dr.-Ing.
Advisor: Samuel Torre Vinuesa, Jerome Mazerier
Submission Date: Submission date

I confirm that this master's thesis is my own work and I have documented all sources and material used.

Munich, Submission date

Jih-Wei Liang

Acknowledgments

Abstract

Contents

Acknowledgments	iv
Abstract	v
1 Introduction	1
2 Background	3
2.1 System Monitoring	3
2.1.1 Black Box Monitoring	3
3 Overview	5
3.1 Section	5
3.1.1 Subsection	5
4 Design	6
4.1 Section	6
4.1.1 Subsection	6
5 Implementation	7
5.1 Section	7
5.1.1 Subsection	7
6 Evaluation	8
6.1 Section	8
6.1.1 Subsection	8
7 Related Work	9
7.1 Using Prometheus for Black Box Monitoring	9
7.2 Datadog in the Context of Black Box Monitoring Systems	10
8 Summary and Conclusion	12
8.1 Section	12
8.1.1 Subsection	12

Contents

9 Future Work	13
9.1 Section	13
9.1.1 Subsection	13
Abbreviations	14
List of Figures	15
List of Tables	16
Bibliography	17

1 Introduction

Black box monitoring is a system analysis technique used in software testing and network monitoring. It assesses system functionality based on its output in response to certain inputs, without considering its internal mechanisms, thereby treating the system as a “black box” [BW96] [Hie06].

Uniping designed to enable black box monitoring for applications of Open Back End (OBE) and Service Integration (SI) in Amadeus Cloud Service (ACS) was a deprecated project in Amadeus. Recently, demands for black box monitoring in ACS have increased, leading to the reboot of this project. Therefore, the goal for Amadeus is to enhance Uniping in ACS for distributed and reliable Black Box Monitoring and contribute Uniping to the Open Source Community as an innovative solution in the realm of system monitoring.

Uniping is composed of three main components: the Operator, the Scheduler, and the Pinger. The Operator is responsible for handling the Custom Resource Definition (CRD) that defines the service and related information for monitoring. The Scheduler is tasked with scheduling monitoring requests based on the CRD and collecting data from the Pingers. Lastly, the Pinger is a custom worker designed for various monitoring purposes. It sends requests to targets and then returns the necessary data back to the Scheduler.

However, Uniping has several limitations, including a single point of failure, lack of scalability, and limited customizability. Specifically, the Scheduler, which dispatches monitoring requests and collects response data, represents a single point of failure and could potentially become a traffic bottleneck. Furthermore, Uniping can only be deployed to OpenShift Clusters individually, with no common control plane for all Unipings deployed across different clusters. These design and architectural issues pose significant challenges for the rebooted project, especially for other teams within Amadeus who rely on Uniping for reliable black box monitoring of a large number of distributed and cross-cluster services.

Nowadays, there are renowned tools to implement system monitoring such as Prometheus, Datadog, Nagios, and so on. These tools aim to collect a variety of metrics from targets for monitoring computational resources. Generally, they provide an agent to scrape desired data and then wait for the server’s request or actively push these data to the responsible server. Most of them feature great professionals

in the realm of monitoring, offering enterprise-level services in reaction to kinds of requirements [NVP21].

Uniping aims to provide real-time monitoring, logging, and analysis capabilities. These features enable system administrators to quickly detect and diagnose issues, thereby improving system reliability and performance [JD21]. Furthermore, the proposed improvements to Uniping will offer scalability, availability, and failure tolerance through a redesigning and refactoring of the system architecture. Consequently, the objective of the project is to redesign and implement an open-source black box monitoring system that meets current requirements in the context of distributed architecture.

2 Background

2.1 System Monitoring

2.1.1 Black Box Monitoring

Black box monitoring is a well-established approach to system analysis, most commonly applied in software testing and network monitoring. It evaluates a system's functionality based on its responses to given inputs while abstracting away the system's internal workings. This technique metaphorically views the system as a "black box", signifying that the internal mechanisms are not visible or accounted for in the evaluation process [BW96] [Hie06].

One of the key advantages of black box monitoring is its simplicity and ease of implementation [Jai91]. Since it doesn't require in-depth knowledge of the system internals, it can be quickly deployed across various platforms and systems. This quality fosters wide applicability, accommodating diverse programming languages and system architectures. Furthermore, black box monitoring has a unique focus on the user experience. By simulating user actions and recording the system's responses, it offers a valuable measure of system functionality from a user's perspective [JD21]. Thus, it is instrumental in ensuring that the system effectively meets the end-user's requirements. However, while the black box approach has its merits, it is not without challenges. One significant limitation is the potential for limited coverage. Since black box monitoring concentrates on the inputs and outputs, it might overlook internal system issues [KFN99]. If a problem doesn't immediately or directly affect the system output, it may remain undetected, only to potentially cause serious issues down the line. Moreover, black box monitoring could contribute to inefficiency in troubleshooting. A system failure detected by black box testing can be difficult to diagnose and fix due to the lack of visibility into the system internals. Pinpointing the specific area of the system causing the issue may turn into a challenging endeavor [KFN99].

Despite these challenges, the future of black box monitoring looks promising, particularly in the context of our rapidly digitalizing world. Given the rising importance of customer satisfaction and user experience in the digital economy, black box monitoring's ability to capture system performance from a user's viewpoint remains invaluable. The evolution of Artificial Intelligence (AI) and Machine Learning (ML) technologies

could significantly enhance the efficiency of black box monitoring. AI-powered monitoring tools can automate the process of generating test inputs and predicting system outputs, bringing about a higher degree of efficiency and reliability [Ber03]. As the IT landscape becomes more complex with microservices and cloud-based applications, understanding the internals of every service can be overwhelming. In such a scenario, black box monitoring's ability to provide a holistic view of system behavior can be highly beneficial. Moreover, in light of growing privacy regulations and data protection measures, black box monitoring's non-invasive approach aligns with the current trends. Focusing on system outputs rather than internals could potentially minimize privacy or data protection concerns [OA07].

In conclusion, black box monitoring remains a vital tool for system testing and monitoring, despite its inherent challenges. Its future appears increasingly integrated with AI technologies and aligned with user-centric design principles. Nevertheless, black box monitoring should not be considered a stand-alone solution but rather a component of a comprehensive monitoring strategy that incorporates various techniques to effectively manage system performance.

3 Overview

3.1 Section

3.1.1 Subsection

4 Design

4.1 Section

4.1.1 Subsection

5 Implementation

5.1 Section

5.1.1 Subsection

6 Evaluation

6.1 Section

6.1.1 Subsection

7 Related Work

7.1 Using Prometheus for Black Box Monitoring

The Prometheus Blackbox Exporter is a valuable extension of the Prometheus monitoring system, specifically designed for monitoring external services and endpoints. This tool permits users to probe targets via HTTP, TCP, and ICMP protocols, yielding metrics related to availability, latency, SSL certificate expiration, and DNS resolution, among others. Its configuration options allow for the customization of probe settings, including timeout periods, headers, and probe methods. The metrics gathered are presented in a Prometheus-compatible format, enabling seamless integration with existing Prometheus monitoring setups [Pro] [Pro23].

The Prometheus community is renowned for its high-quality software, a variety of plugins, and robust support, which can be leveraged when creating a Blackbox Monitoring system. According to a presentation by Aaron Wiecek titled "Spotlight on Cloud: Using Prometheus for Black Box Monitoring," utilizing a Prometheus Blackbox Exporter based architecture for blackbox monitoring is a practical choice for systems already leveraging Prometheus for monitoring and alerting [ORE]. The following figure depicts the setup:

However, there are a few challenges. Prometheus' default data retention period is limited to fourteen days because its primary function is real-time system status monitoring and alert delivery. It is not built to retain large amounts of data for extended periods. Furthermore, Prometheus is not designed for a distributed setup, increasing complexity when configuring and deploying each Blackbox Exporter separately, especially during inevitable configuration changes (see Figure 2) [Pro]. Another critical drawback is that Prometheus acts as a single point of failure, jeopardizing the availability of the monitoring system and risking data loss during a system failure.

In conclusion, implementing black box monitoring with the Prometheus Blackbox Exporter comes with benefits and drawbacks [ORE]. While it is praised for its simplicity, customization, and extensive community support, it has inherent limitations from its design, constraining its ability to fully exploit the power of a distributed architecture. Data retention for analytics is increasingly vital for ensuring the Service Level Agreement (SLA) of private services. Without centralized control over all Blackbox Exporters, configuration management can become unwieldy, and the single point of failure inher-

ent in the design presents a risk of data loss. Therefore, careful consideration of these trade-offs is essential in the design and decision-making process of the project.

7.2 Datadog in the Context of Black Box Monitoring Systems

In the realm of modern IT infrastructure management, monitoring and analytics platforms play an integral role in ensuring system reliability and performance. Among the existing solutions, Datadog has emerged as a prominent player, known for its robust functionalities and capabilities [Datd]. In the context of this project, which aims to design a dedicated and distributed black box monitoring system, understanding Datadog's approach and technology is pivotal.

Datadog offers a unified view of an organization's IT infrastructure by collecting and analyzing data from multiple sources. This approach enables real-time monitoring and troubleshooting, as well as optimization of applications and infrastructure. A standout feature of Datadog is its Synthetic Monitoring, which is particularly applicable for black box monitoring [Datd] [Data]. This feature simulates requests and actions, thereby offering insights into the performance of APIs across various network layers, from backend to frontend [Data] [Datb].

There are two distinct elements of Datadog's Synthetic Monitoring that are of particular interest: API Tests and Browser Tests [Data] [Datb]. API Tests actively probe target services, gathering their statuses and metrics. These tests offer a snapshot of the availability of the monitored services, a crucial factor in any black box monitoring system. In Datadog's implementation, API Tests form the bulk of the monitoring process, thereby highlighting their importance. Browser Tests, on the other hand, execute defined scenarios on target services using a chosen browser. This essentially reproduces the actions of a user, providing experiential feedback about the services. This user-centric approach complements the more technical data gathered by the API Tests, contributing to a more holistic view of the system's health.

When it comes to monitoring availability, it is crucial to decide the position of source endpoints that proceed with monitoring [Datc]. Datadog provides the feature to customize private locations to execute Synthetic Monitoring. By installing Docker containers as private endpoints in desired locations, Synthetic Monitoring, including both API Tests and Browser Tests, can employ these private endpoints to compose its line of departure.

While Datadog's synthetic monitoring approach presents a well-rounded solution, it operates primarily within a centralized model in terms of the agents' side. To elaborate, monitoring targets from a private location requires an installed agent in the same cluster, and the agent could be a single point of failure.

This brings us to the primary focus of our thesis: to explore the feasibility and benefits of a distributed black box monitoring system. As businesses scale and become increasingly distributed, there may be unique advantages to employing a monitoring system that mirrors this distributed structure. By building on the strengths of Datadog's approach, while addressing its potential limitations in distributed environments, we aim to advance the field of black box monitoring.

8 Summary and Conclusion

8.1 Section

8.1.1 Subsection

9 Future Work

9.1 Section

9.1.1 Subsection

Abbreviations

ACS Amadeus Cloud Service

AI Artificial Intelligence

CRD Custom Resource Definition

ML Machine Learning

OBE Open Back End

SI Service Integration

SLA Service Level Agreement

List of Figures

List of Tables

Bibliography

- [Ber03] A. Bertolino. "Software testing research and practice." In: *Proceedings of the abstract state machines 10th international conference on Advances in theory and practice*. ASM'03. Berlin, Heidelberg: Springer-Verlag, Mar. 3, 2003, pp. 1–21. ISBN: 978-3-540-00624-4.
- [BW96] B. Beizer and J. Wiley. "Black Box Testing: Techniques for Functional Testing of Software and Systems." In: *IEEE Software* 13.5 (Sept. 1996). Conference Name: IEEE Software, pp. 98–. ISSN: 1937-4194. DOI: 10.1109/MS.1996.536464.
- [Data] Datadog. *API Tests*. Datadog Infrastructure and Application Monitoring. URL: https://docs.datadoghq.com/synthetics/api_tests/ (visited on 05/21/2023).
- [Datb] Datadog. *Browser Tests*. Datadog Infrastructure and Application Monitoring. URL: https://docs.datadoghq.com/synthetics/browser_tests/ (visited on 05/21/2023).
- [Datc] Datadog. *Run Synthetic Tests from Private Locations*. Datadog Infrastructure and Application Monitoring. URL: https://docs.datadoghq.com/synthetics/private_locations/ (visited on 05/21/2023).
- [Datd] Datadog. *Synthetic Monitoring*. Datadog Infrastructure and Application Monitoring. URL: <https://docs.datadoghq.com/synthetics/> (visited on 05/21/2023).
- [Hie06] R. M. Hierons. "Software Testing Foundations: A Study Guide for the Certified Tester Exam. By Andreas Spillner, Tilo Linz and Hans Schaefer. Published by dpunkt.verlag, Heidelberg, Germany, 2006. ISBN: 3-89864-363-8, pp 266: Book Reviews." In: *Softw. Test. Verif. Reliab.* 16.4 (Dec. 1, 2006), pp. 289–290. ISSN: 0960-0833.
- [Jai91] R. Jain. *The Art of Computer Systems Performance Analysis: Techniques For Experimental Design, Measurement, Simulation, and Modeling*, NY: Wiley. Apr. 1, 1991.
- [JD21] P. Jorgensen and B. DeVries. *Software Testing: A Craftsman's Approach*. May 31, 2021. ISBN: 978-1-00-316844-7. DOI: 10.1201/9781003168447.

- [KFN99] C. Kaner, J. L. Falk, and H. Q. Nguyen. *Testing Computer Software, Second Edition*. 2nd. USA: John Wiley & Sons, Inc., Mar. 1999. 480 pp. ISBN: 978-0-471-35846-6.
- [NVP21] F. Neves, R. Vilaça, and J. Pereira. “Detailed black-box monitoring of distributed systems.” In: *ACM SIGAPP Applied Computing Review* 21 (Mar. 1, 2021), pp. 24–36. DOI: 10.1145/3477133.3477135.
- [OA07] P. N. Otto and A. I. Anton. “Addressing Legal Requirements in Requirements Engineering.” In: *15th IEEE International Requirements Engineering Conference (RE 2007)*. 15th IEEE International Requirements Engineering Conference (RE 2007). ISSN: 2332-6441. Oct. 2007, pp. 5–14. DOI: 10.1109/RE.2007.65.
- [ORe] O’Reilly. *Spotlight on Cloud: Using Prometheus for Black Box Monitoring with Aaron Wiczorek*. O’Reilly Online Learning. URL: <https://learning.oreilly.com/videos/spotlight-on-cloud/0636920360216/0636920360216-video328883/> (visited on 05/21/2023).
- [Pro] Prometheus. *Overview | Prometheus*. URL: <https://prometheus.io/docs/introduction/overview/> (visited on 05/26/2023).
- [Pro23] Prometheus. *Blackbox exporter*. original-date: 2015-09-05T14:45:00Z. May 26, 2023.